Two estimation methods on the same texts

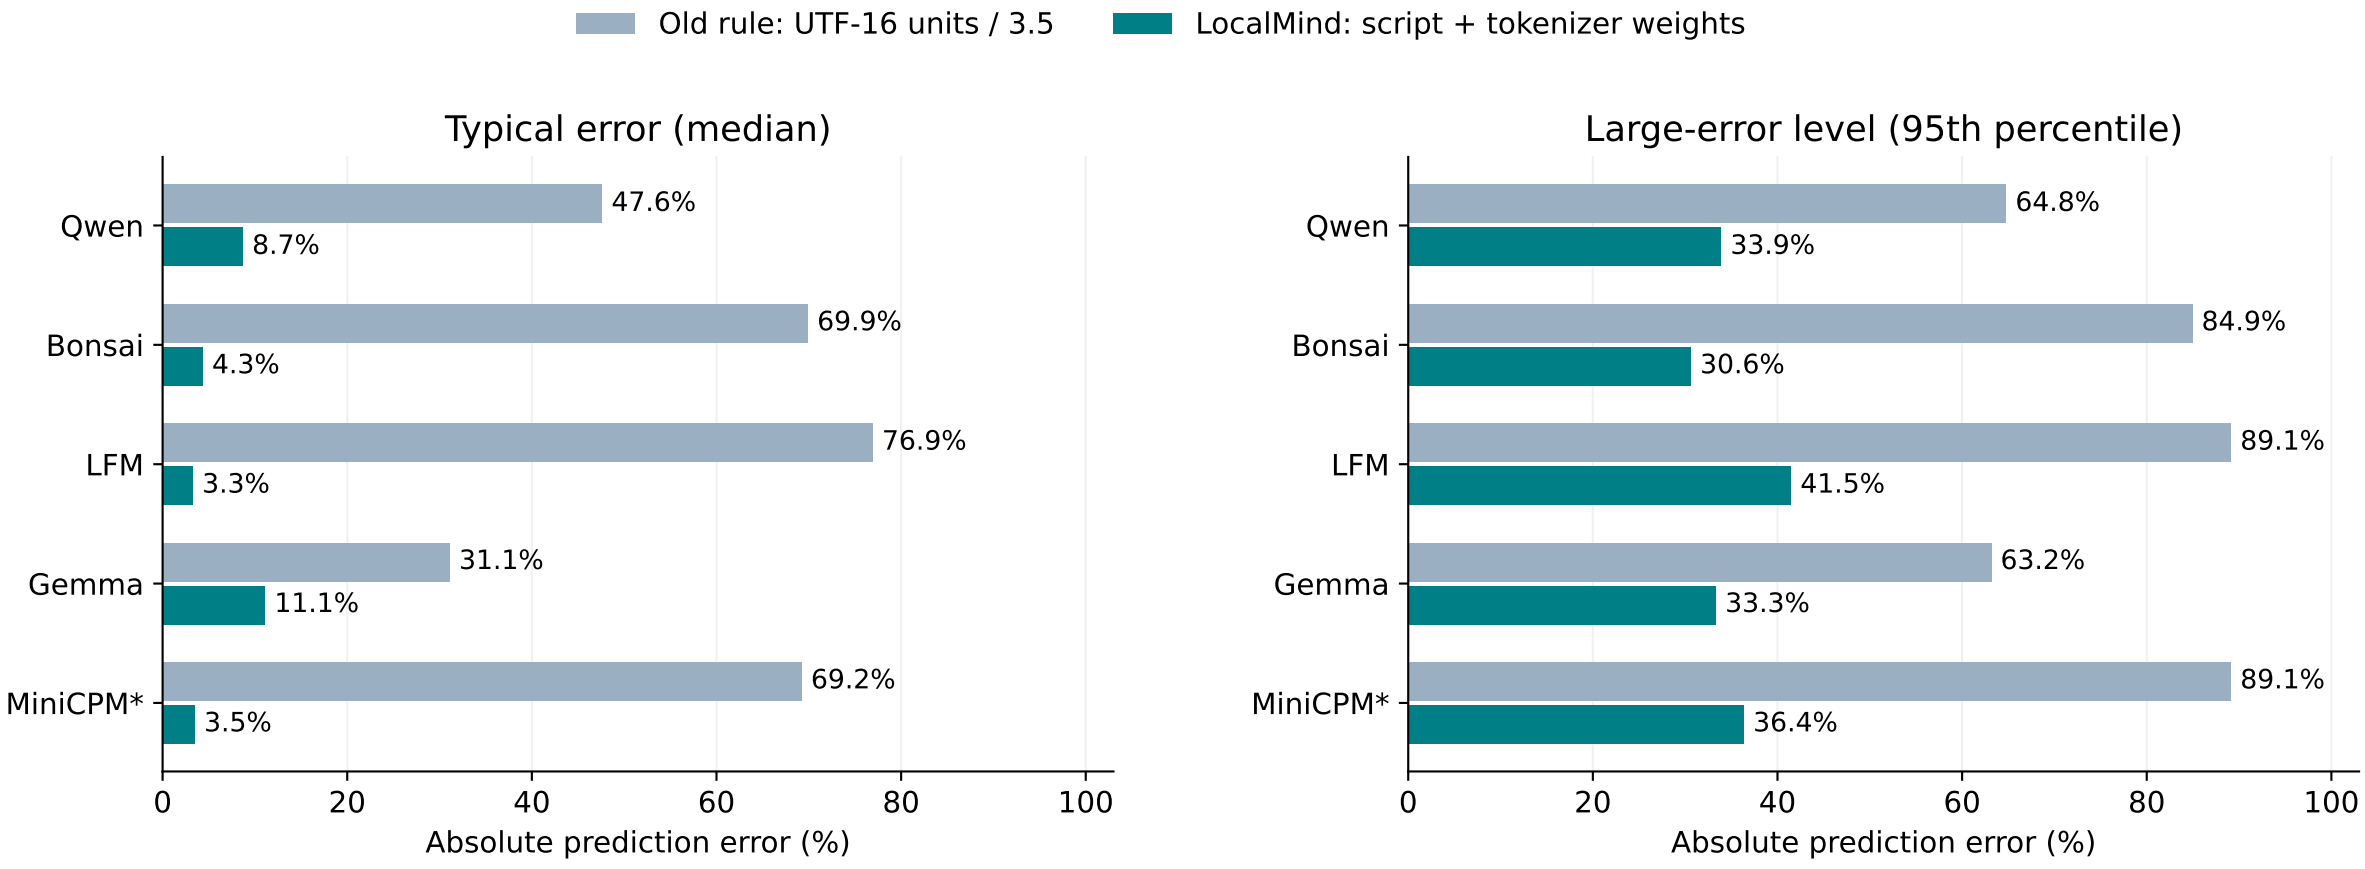

1,800 texts per tokenizer • 18 language varieties • FLORES+ dev pilot

Smaller is better. Absolute error combines under- and overestimation; the 95th percentile is not a maximum.

*MiniCPM: HF JSON tokenizer. Neither fitting nor a browser benchmark. Same 100 aligned sentence IDs across languages.