Tokenizer differences depend on the language

Same text in each pair • Dot: median ratio • Line: 10th–90th percentile of observed ratios

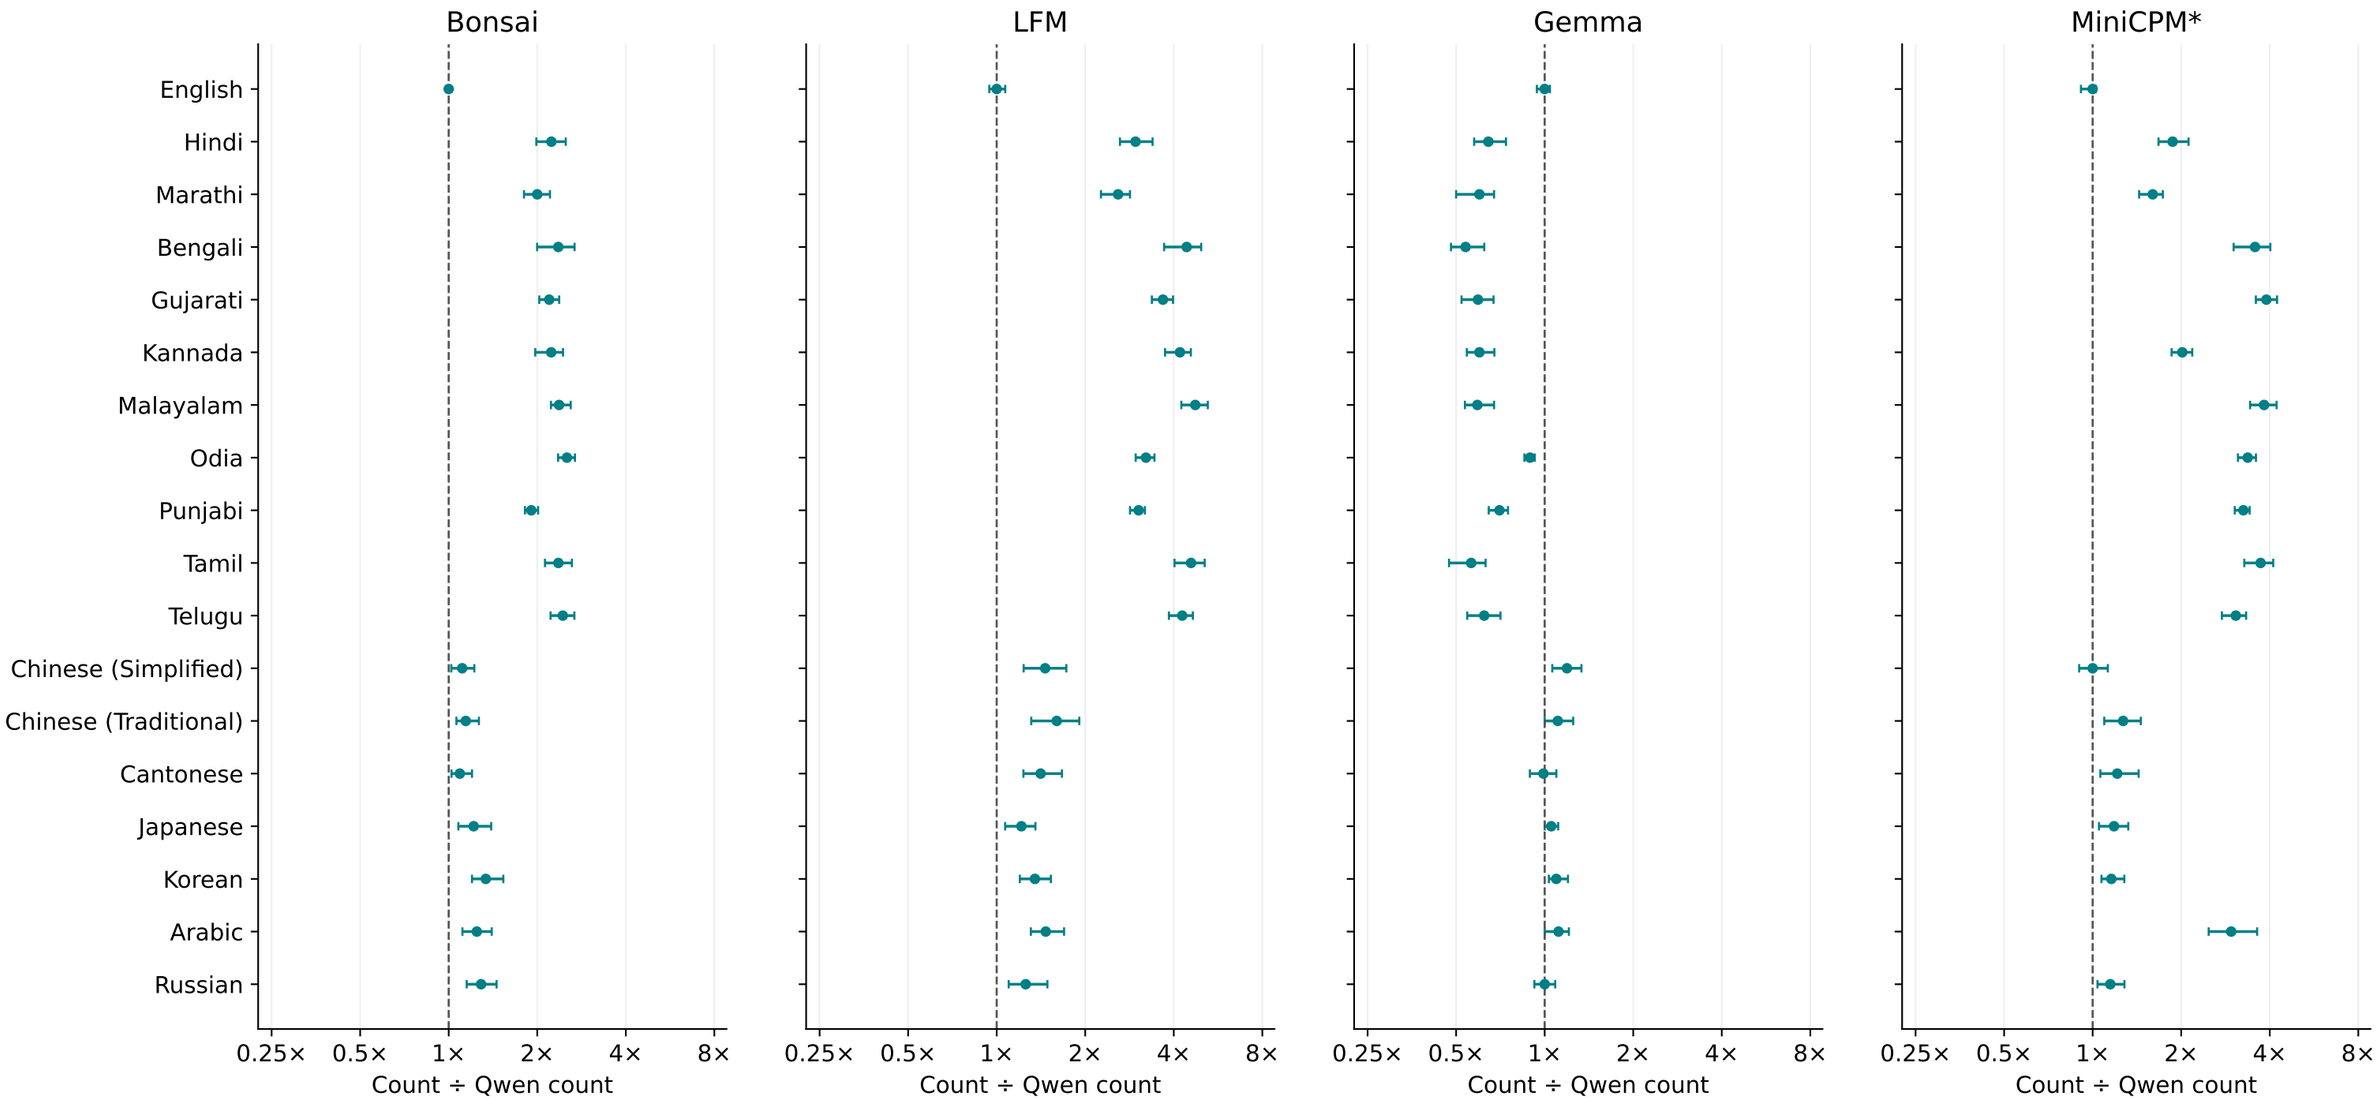

100 paired texts per row. Lines show sample variation, not confidence intervals. Same logarithmic scale in all panels.

Qwen is a fixed reference, not a quality standard. All tokenizer pairs are retained in the numeric results.

FLORES+ dev • Content tokens only • Five BPE profiles • *MiniCPM uses its HF JSON tokenizer.