Same input: token counts relative to Qwen

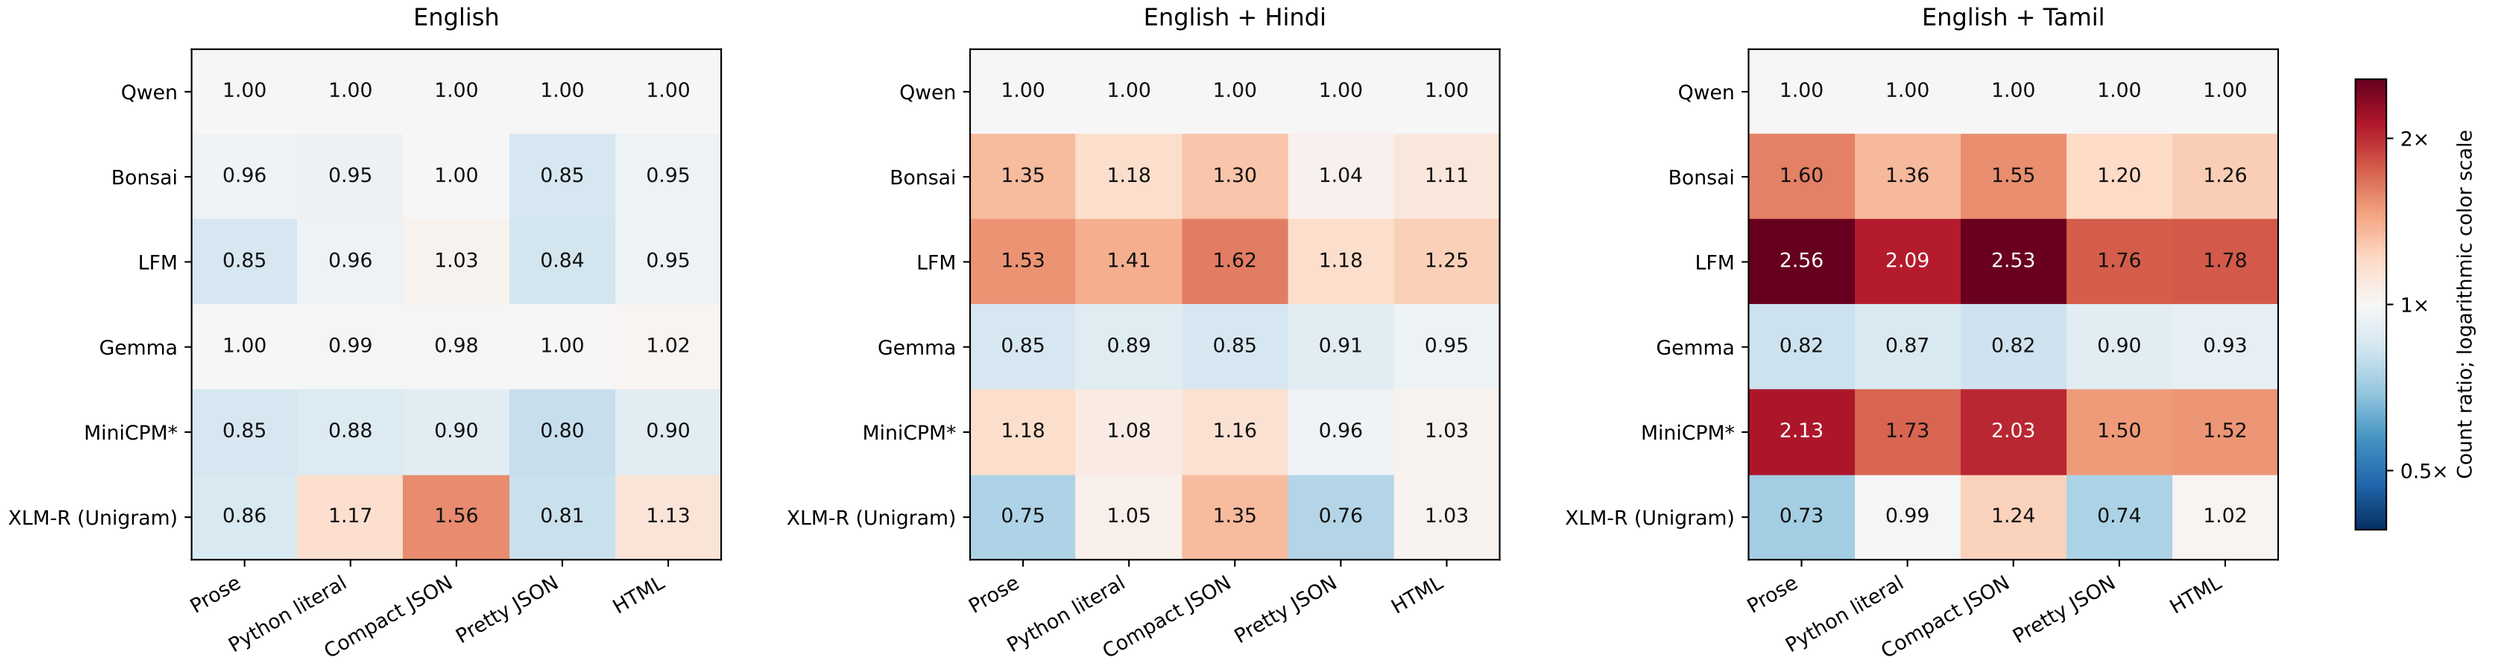

Geometric mean of paired count ratios. 1 = equal; 2 = twice as many. No unknown tokens occurred in primary inputs.

12 authored scenario blocks per cell; 180 primary inputs total. Shared templates, not a natural-workload sample.

XLM-R is one Unigram profile; the other five are BPE. No causal algorithm ranking. MiniCPM uses HF JSON.