Does decoding reproduce the original text exactly?

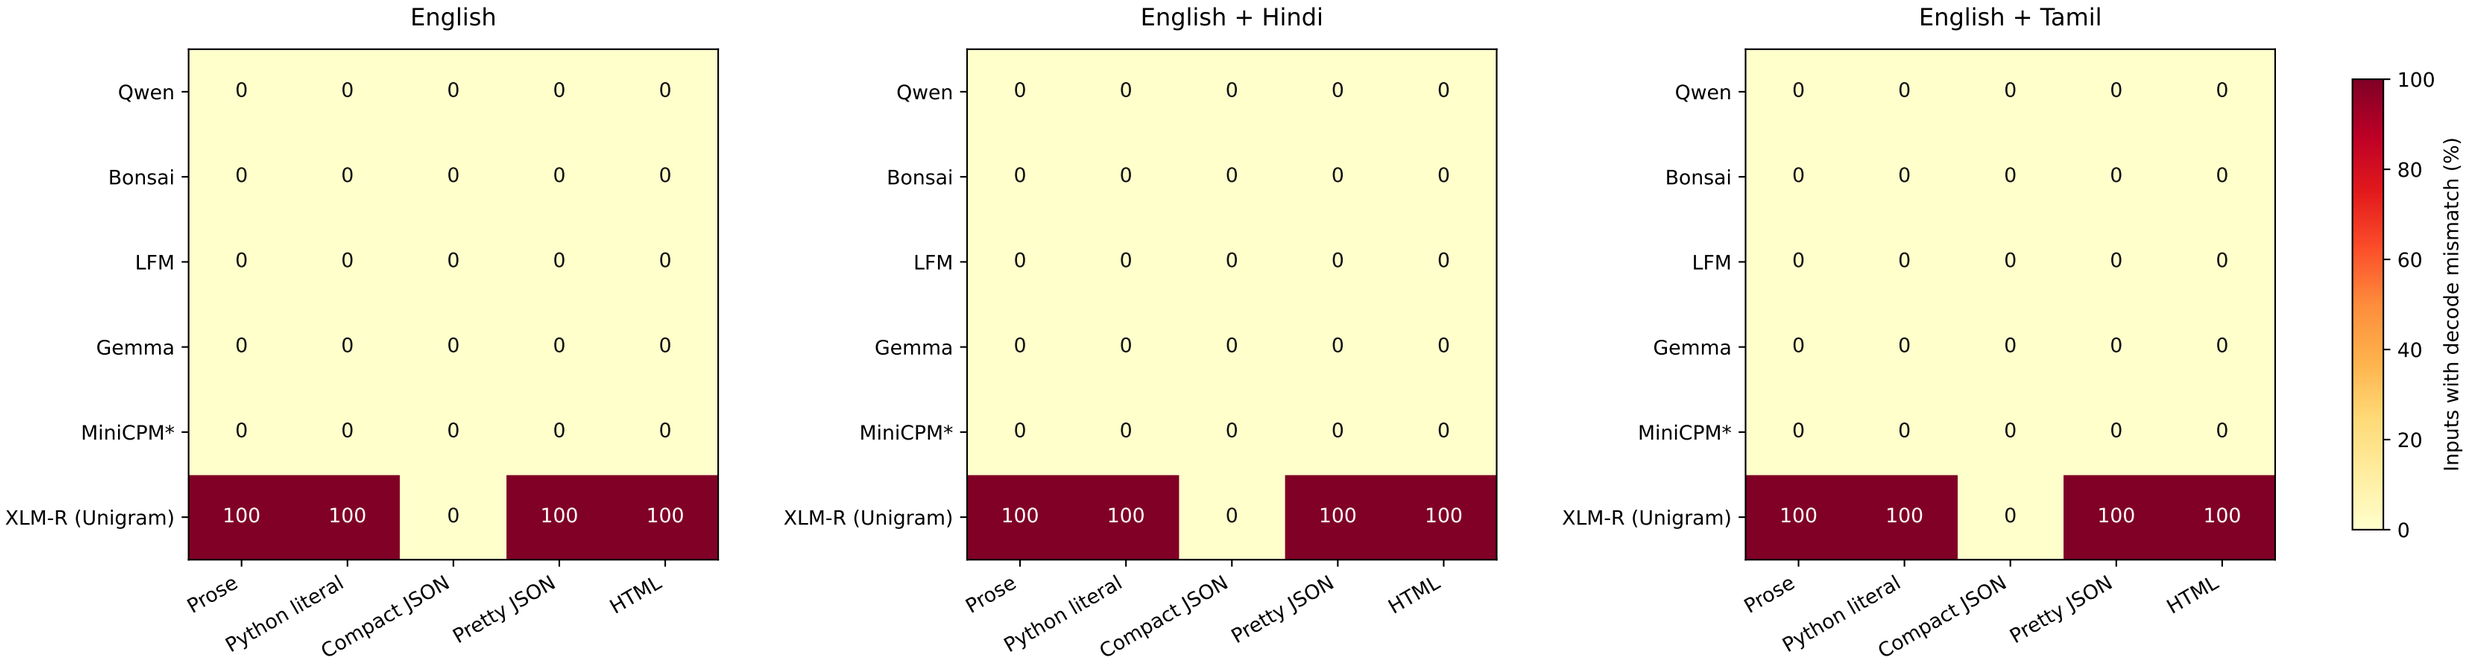

100 = all 12 inputs changed; 0 = all matched exactly. Whitespace/normalization changes are not automatically semantic loss.

12 authored scenario blocks per cell; 180 primary inputs total. Shared templates, not a natural-workload sample.

XLM-R is one Unigram profile; the other five are BPE. No causal algorithm ranking. MiniCPM uses HF JSON.