Different text groups change the relative token counts

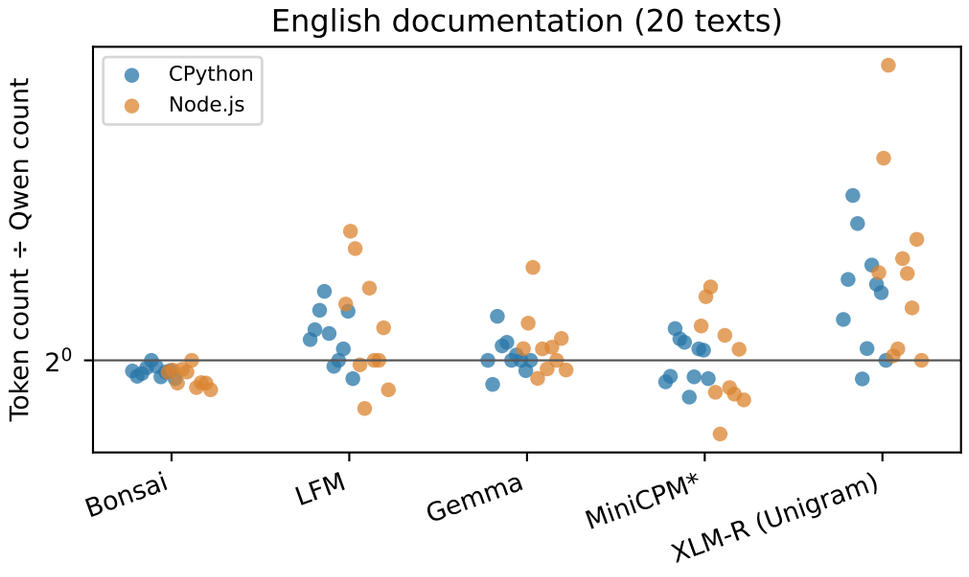

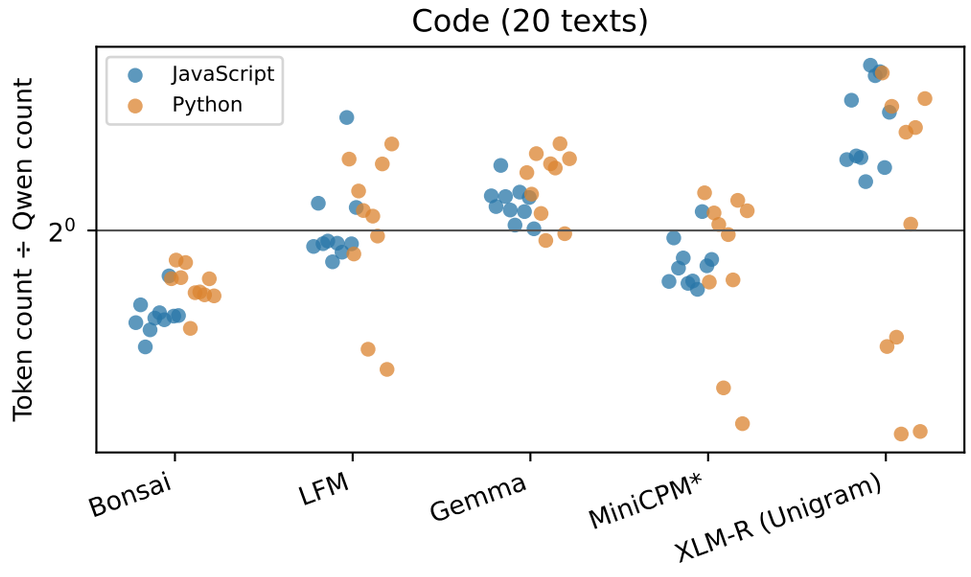

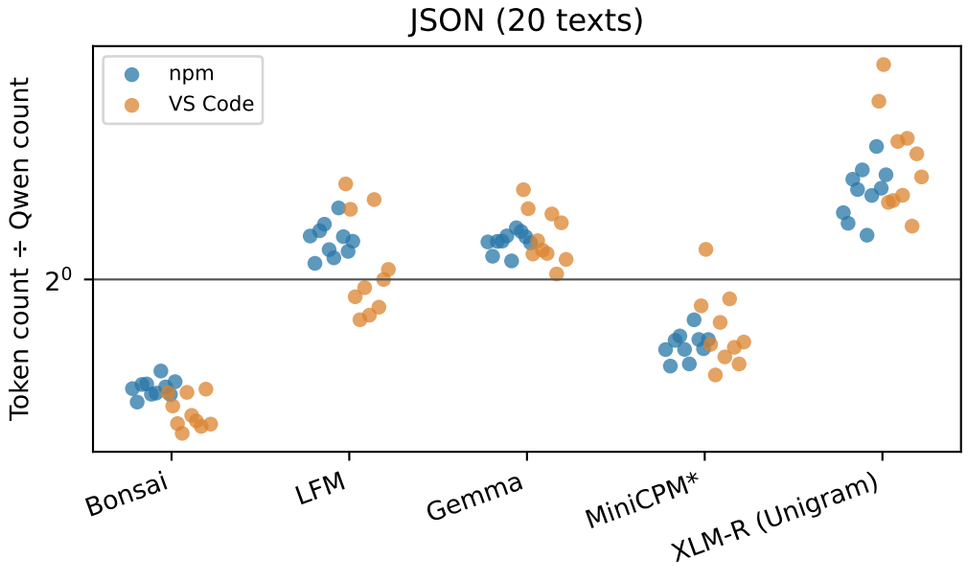

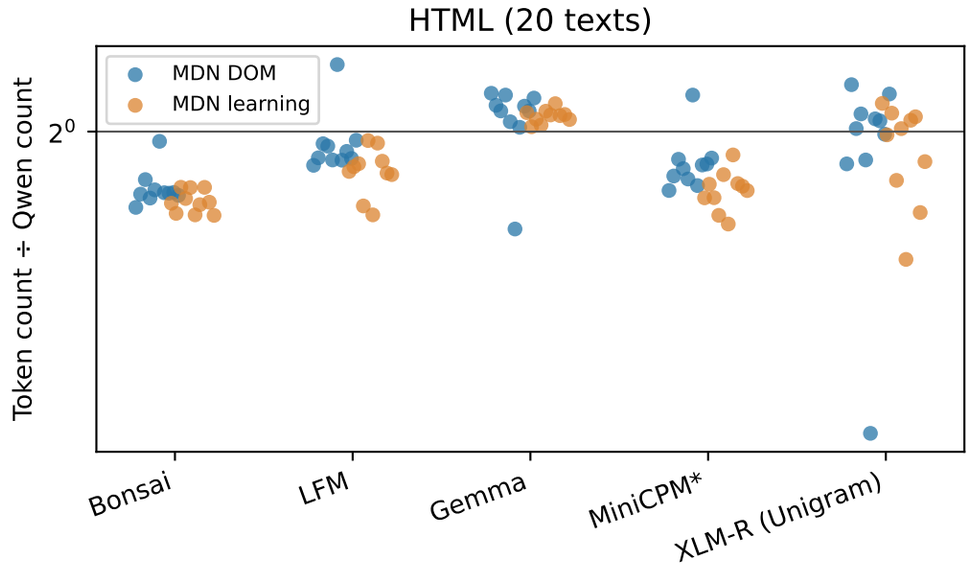

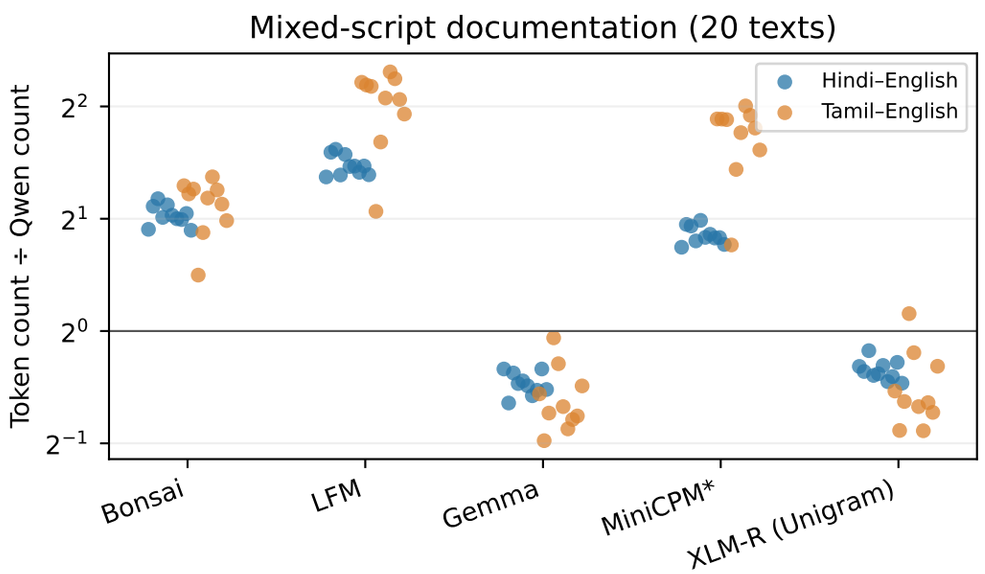

Each dot compares the SAME input.
1 = equal counts; 2 = twice as many.
Colors show source or language.
Qwen is the reference, not a quality standard.

Panels use logarithmic axes with different ranges.
Five BPE profiles + one Unigram profile;
this does not isolate an algorithm effect.

100 published technical inputs; shared projects/templates, automatic translations. MiniCPM*: HF JSON, not verified GGUF parity.